

Name: _____ Student ID: _____

COS80013 Internet Security

Demo 3 (week 3)

You will need:

RedHat Linux 7.3 (VM)

Windows XP Pro (VM)

Windows 95 (VM)

A computer with internet access

In this lab you will experiment with Buffer overflows in the C language.

1. Using Virtual Machine Launcher, start up the COS80013 / **RedHat Linux** with local network VM image.

Alternatively zipped copies are on Cloudstor here:

[Virtual Machines - OneDrive](#)

Log in as *student* (user)

student (password)

Let's write a C program:

Start the editor thus:

pico memtest1.c

Type in the following code:

```
/* memtest1.c */
#include <stdio.h>
#define SIZE 8
void test(int*, char*);
int main() {
    int i = 0;
    char buf[SIZE];
    printf("Type in 5-20 chars into the text buffer\n");
    printf("Watch the value of i \n");
    printf("it will be corrupted when you exceed %i chars\n", SIZE);
    printf("Type \"q\" to exit:\n");
    printf("\t| i posn\t| buf start\t| buf end\t| i value\n");
    do {
        test(&i, buf);
        i++;
    } while(buf[0] != 'q');
    return 0;
}

void test(int *j, char* buf) {
    scanf("%s", buf);
    printf("OK.\t| %u\t| %u\t| %u\t| %d\n", j, buf, &buf[SIZE], *j);
    return;
}
```

Use

Control+O to write to file (followed by ENTER)

Control+X to exit

Name: _____ Student ID: _____

Compile thus:

```
gcc -o memtest1 memtest1.c      (creates executable)
```

When the function is called, **i** is passed by reference. The **test** function can access this variable.

buf, being an array, is passed by reference, so the memory location is shared between the main function and **test**.

```
chmod +x memtest1      (adds execute permissions)
./memtest1              (runs the program)
```

Where in memory (the address) is **i**? What is its value?

	Location (posn)	Comments
i		<i>1</i>
buffer start		<i>Start of user input</i>
buffer end		<i>should be \0</i>
buffer size		
Bytes between buffer start and i		<i>e.g. 3221223812 – 3221223800 = 12 input of more than 12 bytes should start to corrupt memory (i.e. value of i)</i>

Try to find the largest number of chars before a buffer overflow or segfault.
Try to make the value of **i** jump.

When a buffer fills up, it writes forwards or backwards depending on the CPU and operating system.

Type in a few strings smaller than 10 characters – note that the integer **i** is counting correctly.

Play with the program (type stuff into it) and record your observations here. Progressively input more characters into the program and note changes on behaviour.

Test No (i)	Text typed into buf	Number of characters	Behaviour
0	<i>12345</i>	<i>5</i>	<i>Increments i</i>
1			
2			

Name: _____ Student ID: _____

3			
4			
5			<i>i jumps to ...?</i>
6			
7			
8			

What is the distance (in bytes) between **buffer start** and **i**?

$n = (&i - \text{buffer start}) =$

Try inputting that many characters (**n**) to overwrite the variable **i**.
Try it. Does *i* change?

Try inputting $n + 1$ characters to overwrite the variable **i**.
Try it. Does *i* change?

If you type 12 characters and <Enter>, the computer inserts a NULL character ('\0') in the 13th position.

If it crashes, just start it up again and keep typing longer strings. Watch the value of i.

If a string is really long, you will crash the program (Linux will report a Segmentation Fault).

Code can be injected into the window between buffer overflow and segmentation fault, overwriting other parts of the program such as the **return address** and the **EBP**

More background here:

<http://www.tenouk.com/Bufferoverflowc/Bufferoverflow2a.html>

2. Try this program:

```
/* memtest2.c */
#include <stdio.h>
int main()
{
    char first[12];
    char last[12];

    printf("Type in your first name: ");
    gets(first);
}
```

Name: _____ Student ID: _____

```
printf("Type in your last name: ");
gets(last);

printf("Hello %s %s\n", first, last);

return 0;
}
```

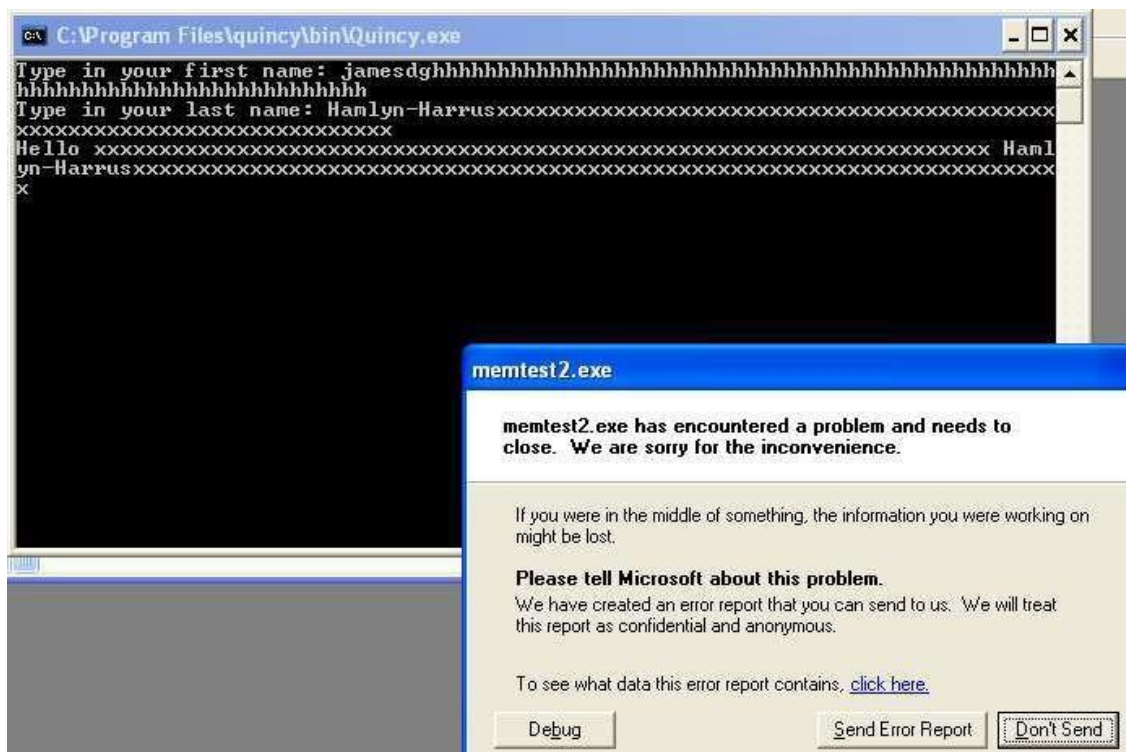
The compiler warning about **gets** will not stop compilation. You can ignore it.

Unlike **scanf**, **gets** lets you type in spaces and other non-printing characters, so it allows an attacker to enter executable code.

If you put in long strings, you can overflow one string with the contents of the other. If they are big enough, you will get a segmentation fault. If the size is just right, you will overwrite a function somewhere else in the program.

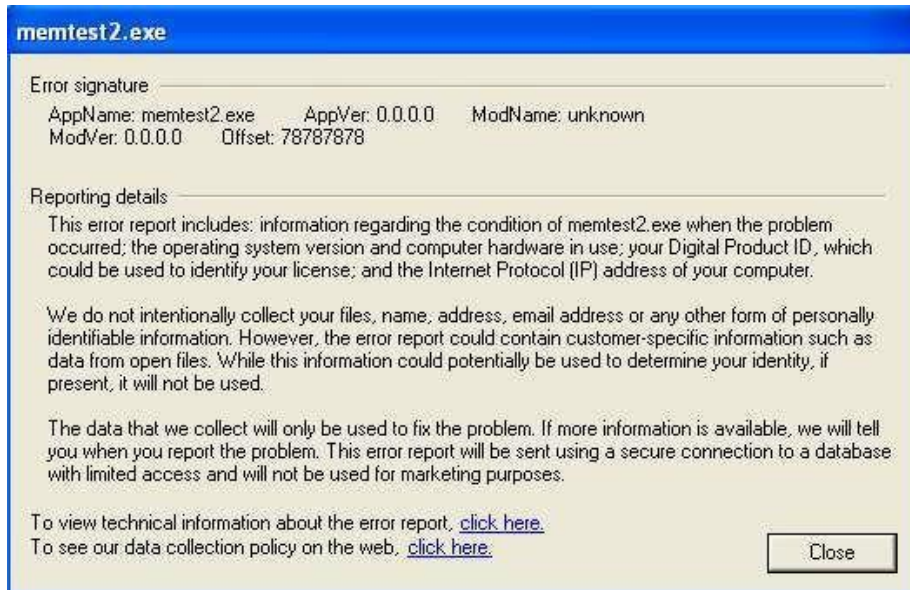
How many characters must you input into the last name to overwrite the first name (try it)?

3. (optional) You can also try these programs in Windows using the Quincy IDE (in the **Windows 95** VM). If a 'segmentation error' occurs, Windows will pop up this message:



Name: _____ Student ID: _____

If you click on [click here](#), you will see the contents of the stack pointer, which contains the hex value of some of the string you typed in.



Note: 78787878 is hex for xxxx.

Hackers can use this feedback message to figure out which bytes in their exploit string will be copied into the stack pointer. Changing these critical bytes to the location of some (of their) executable code allows the exploit to run their code. With Linux, it's a bit harder, but not much.

Name: _____ Student ID: _____

4. Re-write **memtest1.c** or **memtest2.c** to prevent excess characters being pasted into the char arrays:

HINT:

replace **scanf("%s")** with **scanf("%10s")**

replace **gets** with **fgets**

usage: `char * fgets (char * str, int num, FILE * stream);`
e.g. **fgets(first, 12, stdin)**

Unlike gets, fgets leaves characters in the keyboard buffer.

In windows you can clean up the input stream with

fflush(stdin);

In Linux you can do the same with

char ch; /*declare once only*/

while ((ch = getchar()) != '\n' && ch != EOF);

example code in **safegets.c**

on Blackboard.

5. Using Google, look up the various "safe" string manipulating functions available in C. (e.g. strncpy, strncat)

What does the *n* do?

How is value of *n* determined?

6. Start up the COS80013 / **Windows XP Pro with local network** VM image

Surf to <http://www.server.com/remote>

Read the file: **readmeb.txt**

Follow the instructions – you will need to log in to the Linux machine to start the server, and then access it from the XP machine using the **passwordclient**.

Name: _____ Student ID: _____

7. On the host PC, surf to <http://www.vividmachines.com/shellcode/shellcode.html> and read up on injecting shellcode into Linux and Windows processes.

What does `.GetProcAddress()` do?

8. Using Google, look for tools and programs for checking a C program for potential buffer overflow vulnerabilities and memory leak testers. **What are the names of some?**

9. Shut down all VMWare images and browsers and log off.

End of Lab

